


Local Search

Local Search and Optimization

- Optimization: To find a state that optimizes (minimizes or maximizes) an **objective function**.
- Here (sections 4.1-4.2 in the textbook) we use searching for the purpose of optimization.
- CSPs (e.g., 8-queens) can be solved using an objective function that represents the degree of satisfaction (or conflict) to the constraints.
- Exploitation-vs-Exploration is still an important issue in local search methods.

State-Space Landscape



- An example of an objective function to be maximized.
- The objective function is a function of the states.
- It is useful to think of the objective function as a surface.

Hill-Climbing Search

- Also called "greedy local search".
- From the current state, compute the objective function for all of its immediate successors.
- Always move to the successor state that optimizes the objective function (steepest ascend / descend).
- Simple and efficient.
- Easy implementation: no queue, no tree, ...
- Can get stuck at suboptimal states:
 - Local maximum / minimum
 - Ridge
 - Plateau (flat regions)

Hill-Climbing Search: Example

An illustration using the 8-queens problem:

- A state has 8 queens, one on each column.
- Each action can move a queen within its column.

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

Which queen to move next,
and to which cell?

- Success rate: 14% (average ≈ 4 steps).
- Probability of getting stuck: 86% (average ≈ 3 steps).

Variants of Hill-Climbing Search

To improve the success rate or efficiency:

- Stochastic hill climbing: Choose probabilistically among several moves)
- First-choice hill climbing: Choose the first move found to be better than the current state (when it is too costly to evaluate all successors)
- Random sideways moves: Allows moves to states with the same objective function
 - Escape from shoulders
 - Need to avoid infinite loops (can limit the number of allowed consecutive sideways moves)
- Random restart

Simulated Annealing

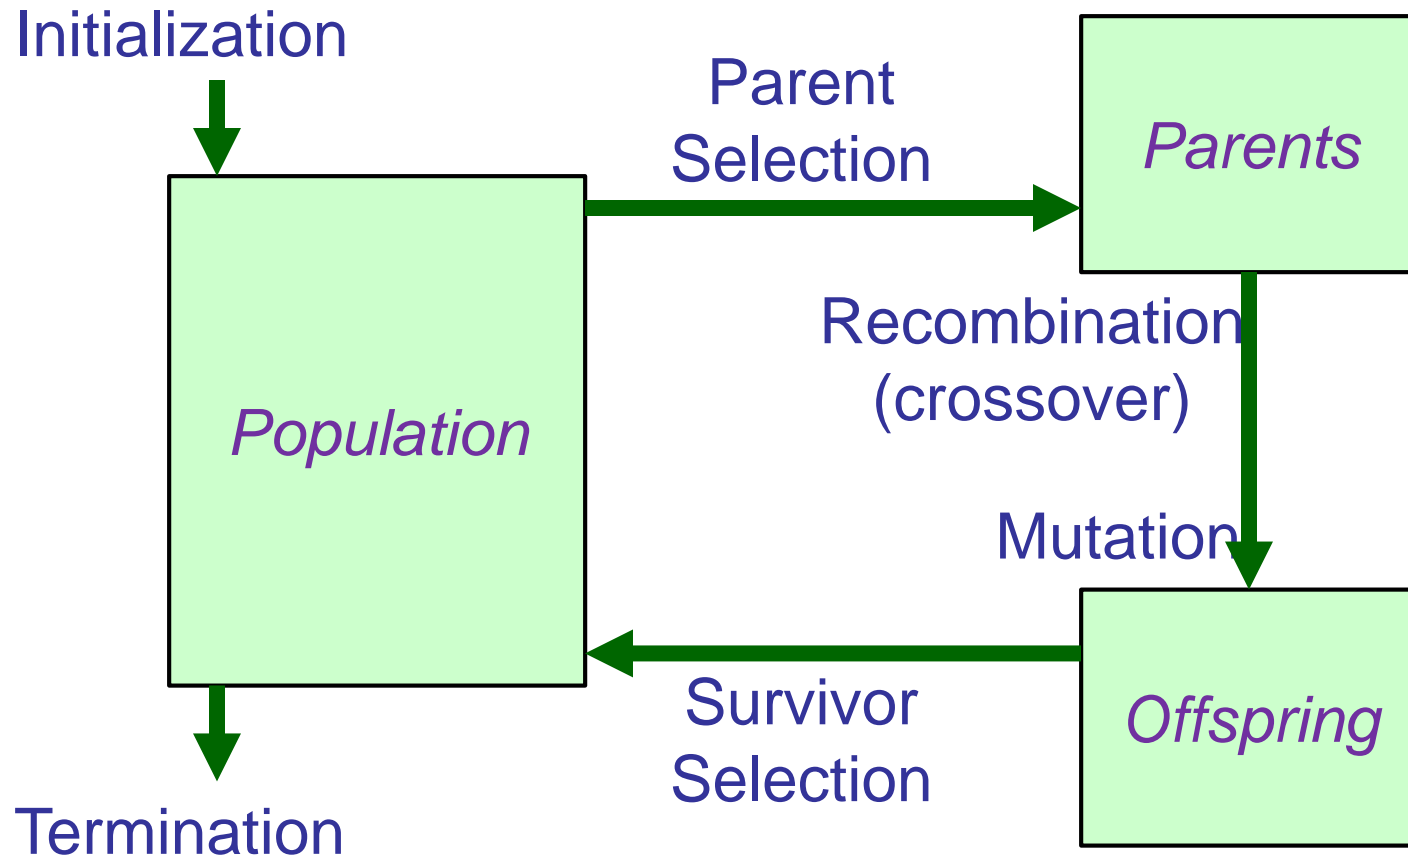
Difference from hill-climbing:

- Rules of selecting a successor:
 - Randomly pick a successor
 - If it is a better one, select it
 - Otherwise, still select it with a probability
- The probability of selecting a worse state is regulated by a variable T (temperature); higher T allows larger probability.
- This leads to opportunities for escaping from local optimums.
- Temperature is gradually reduced to zero (to facilitate convergence).

Genetic Algorithms

- A state is represented as a string. (e.g., the vertical positions of the 8 queens). This is the gene of the state.
- Multiple simultaneous searches by a **population** of individuals (each individual having its own state).
- **Fitness function**: For evaluating the goodness of an individual.
- Each step of the search is a **generation**. Things that can happen in each generation:
 - **Crossover** (two individuals exchange part of their genes)
 - **Mutation** (local move of an individual with a small probability)
 - Selection: Removal of worse individuals and addition of newly generated ones.


How Genetic Algorithms Work



Genetic Algorithms: Example



An illustration of crossover:



Local Search in Continuous Spaces

- Some ways to handle a continuous world:
 - Discretize it!
 - Steepest ascend / descent (a form hill-climbing).
 - Find local maximums/minimums analytically or numerically by setting the gradient of the objective function to zero. (This is not searching.)

Meta-heuristics

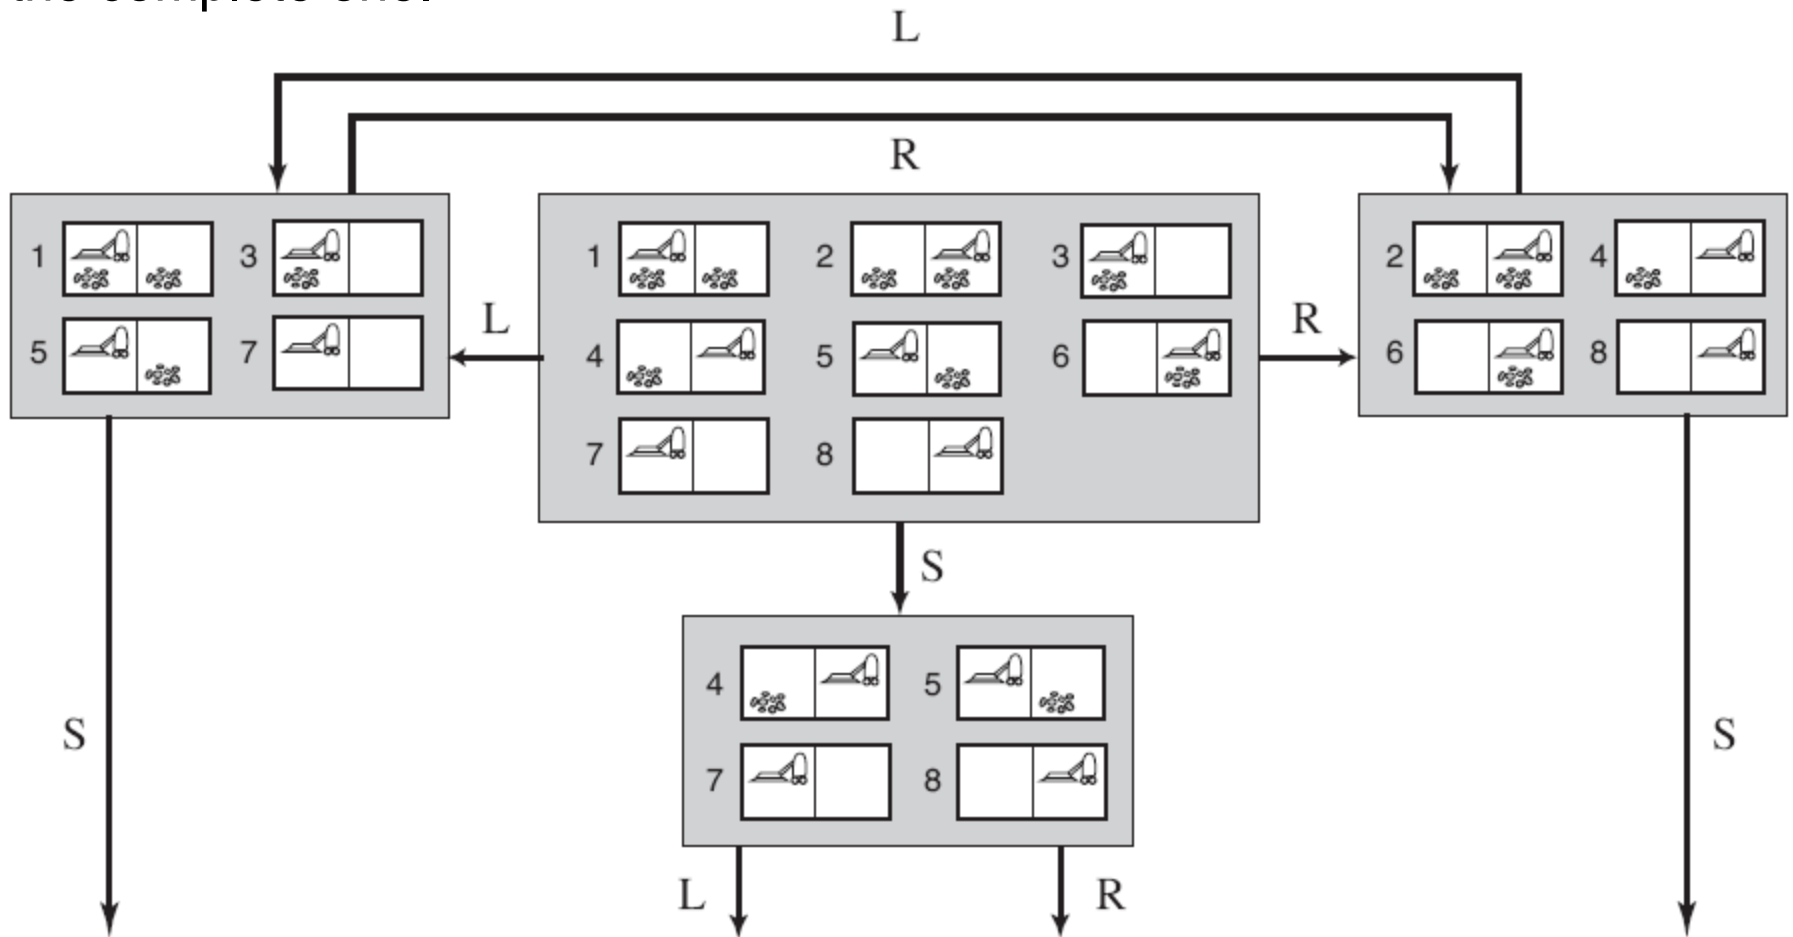
- **Meta-heuristics** are heuristic methods (not functions) for optimization to improve the efficiency and/or likelihood of finding good solutions.
- They are designed to be problem-independent.
- Simulated annealing and genetic algorithms are two representative and widely used meta-heuristics.
- Many are nature-inspired.
- Some other important ones:
 - Particle swarm optimization (PSO)
 - Ant colony optimization (ACO)
 - Artificial immune systems
 - ...

Searching with No Observations

- Known: The state space.
- Unknown: The current state.
- Sensorless (comformant) problem: The agent has no percept to determine the state.
 - **Belief state**: Each belief state contains the possible real states.
 - An action from a belief state results in a new belief state containing the real states that can result from applying the action to any of the real states in the original belief state.
 - Solution can be found by searching in the belief-state space. (A goal state here is one that contains only goal states in the underlying state space.)

Searching with No Observations

Partial belief state space of the vacuum-cleaner world. See the textbook for the complete one.



A solution: [*Right, Suck, Left, Suck*]

Searching with Partial Observations


- Some percept is available.
- The percept received after an action further limits the possible states.
- The solution is a conditional plan (depending on the updated belief state) instead of a sequence of actions.
- Example for the local-sensing vacuum-cleaner world:
 - The percept includes the location and local dirt condition.
 - Initial percept is $[A, \textit{Dirty}]$.
 - Solution: $[Suck, Right, \textbf{if } Bstate = \{6\} \textbf{ then } Suck]$




Here the solution only involves belief states, not actual states.

Searching with Partial Observations

- Example: Localization with local sensing (map is known).
- Percept: Whether there are obstacles in the four (E,W,N,S) directions.
- Action/percept sequence: $NSW \rightarrow Right \rightarrow NS$



(a) Possible locations of robot after $E_1 = NSW$



(b) Possible locations of robot After $E_1 = NSW, E_2 = NS$